

Spark vs. Ray:

The Data Engineering Duel

DA5402 – Assignment 8

Name: Nitesh Kumar Shah
Roll Number: ID25M806
Hardware Partner: Nikesh Kumar Mandal (ID25M805)
Course: DA5402 – MLOps (Jan 2026)
Task: Benchmark Apache Spark against Ray on a 2-node cluster

Stack Summary	
Component	Details
Spark	Apache Spark 3.5.1 (PySpark)
Ray	Ray 2.9.3 with Ray Data
Python	3.11.15
JVM	OpenJDK 17.0.14 (Homebrew)
Dataset	NYC TLC Yellow Taxi (Jan–Jun 2023, 19.5M rows, ~2 GB)
Storage	Apache Parquet + CSV (zone lookup)
Cluster	2 nodes over 802.11ac LAN
Master/Head	MacBook Air M1, 8 cores, 8 GB RAM (192.168.0.134)
Worker	MacBook Air M4, 8 cores, 16 GB RAM (192.168.0.209)

1 Introduction

This assignment investigates the trade-offs between two leading distributed compute frameworks: **Apache Spark** (the JVM-based incumbent of big-data ETL) and **Ray** (the Python-native challenger from the AI infrastructure space). The objective is to implement an identical preprocessing pipeline in both, deploy each on a real two-node cluster, and benchmark performance honestly – with verifiable evidence at every step.

The pipeline targets the much-discussed “Python UDF tax” that Spark allegedly suffers due to its JVM core, against Ray’s native-Python execution model. The dataset is the NYC TLC Yellow Taxi 2023 (Jan–Jun, 19.5M rows, ~2 GB).

Team note: per the assignment, each student worked individually. The 2-node cluster requirement was satisfied by sharing hardware – my code was deployed across my MacBook Air M1 and my partner’s MacBook Air M4 (Nikesh Kumar Mandal, ID25M805). Each of us submits our own implementation and report.

2 Infrastructure Setup

2.1 Cluster Topology

Property	Node 1 (Master/Head)	Node 2 (Worker)
Machine	MacBook Air M1	MacBook Air M4
Cores	8 (4P + 4E)	8
RAM	8 GB	16 GB
LAN IP	192.168.0.134	192.168.0.209
Spark role	Master + Worker	Worker
Ray role	Head + Worker	Worker

Table 1: Cluster hardware inventory. Total capacity: 16 cores, 24 GB RAM.

2.2 Spark Cluster

Spark was deployed using the standalone scheduler. The full Spark 3.5.1 binary distribution was installed on both machines (the `pip install pyspark` package omits the `sbin/` cluster scripts). Master started on the M1; workers on each machine pointed to `spark://192.168.0.134:7077`. The master UI confirmed 2 ALIVE workers (Figure 1).




Figure 1: Spark Master UI showing 2 ALIVE workers (one per Mac) registered to spark://192.168.0.134:7077. Total: 8 cores, 8 GiB.

2.3 Ray Cluster

Ray required two non-default flags: RAY_ENABLE_WINDOWS_OR_OSX_CLUSTER=1 (multi-node Ray on macOS is officially experimental), and -object-store-memory=1500000000 to cap the plasma store given the M1's modest RAM. The dashboard confirmed both nodes ALIVE (Figure 2).




Figure 2: Ray Dashboard: TOTALx2, ALIVEx2. Head at 192.168.0.134 ran at 60% CPU and worker at 192.168.0.209 at 24.7% CPU during live capture.

3 Pipeline Description

The same five-stage pipeline runs in both frameworks.

Stage	Operation
Ingest	Read 6 monthly parquet files (19.5M trips) plus zone CSV. Schema normalisation handles type drift across months (e.g., <code>VendorID</code> INT32 vs INT64; <code>airport_fee</code> casing).
Cleanse	Drop nulls, deduplicate on natural key, filter zero-distance / zero-passenger trips.
Join	Broadcast-join cleaned trips against the 265-row <code>taxi_zones</code> table on <code>PULocationID</code> .
UDF	Apply <code>avg_speed_mph(distance, pickup, dropoff)</code> to every row – the cornerstone of the benchmark.
Export	Write the enriched dataset back to disk as compressed Parquet.

Table 2: The five-stage pipeline – identical logic in both frameworks.

3.1 The Python UDF

Listing 1: Spark: `pandas_udf` for vectorized batch execution

```
@F.pandas_udf(DoubleType())
def avg_speed_mph_udf(distance, pickup, dropoff):
    duration_hr = (dropoff - pickup).dt.total_seconds() / 3600.0
    speed = distance / duration_hr
    return speed.where((duration_hr > 0) & (distance > 0))
```

Listing 2: Ray: `map_batches` with the same vectorized logic

```
def avg_speed_mph_batch(batch):
    duration_hr = (batch["tpep_dropoff_datetime"] -
                  batch["tpep_pickup_datetime"]).dt.total_seconds() / 3600.0
    speed = np.where((duration_hr > 0) & (batch["trip_distance"] > 0),
                    batch["trip_distance"] / duration_hr.replace(0, np.nan),
                    np.nan)
    batch["avg_speed_mph"] = speed
    return batch
```

The arithmetic is identical. The difference: Spark serialises Arrow batches across the JVM↔Python boundary; Ray executes natively in the Python worker process holding the data.

4 Cluster Orchestration Evidence

Beyond the master/dashboard screenshots, two further captures confirm work was *actually distributed* across both Macs.




Figure 3: Spark Executors tab during the distributed run. Three entries: **driver** on .134, executor 0 on .134 (M1) doing 151 tasks, executor 1 on .209 (M4) doing 55 tasks. Both Macs did real work.




Figure 4: Spark master UI after the distributed application finished, showing state **FINISHED** with 8 cores allocated and 3 GiB per executor.

5 Performance Results

5.1 Headline Numbers

Stage	Spark single	Spark 2-node	Ray single	Ray 2-node
1. Ingest	3.62	7.40	28.0	89.60
2. Cleanse	16.36	20.71	38.0	96.13
3a. Join	17.04	16.61	50.0	78.62
3b. UDF	33.65	34.46	70.0	86.05
4. Export	123.53	35.79	70.0	88.60
Total (s)	196.16	116.91	258.99	439.08
Final rows	18,271,535	18,271,535	18,272,001	18,272,001

Table 3: Per-stage execution times (seconds) across all four configurations.

Headline: *Spark beat Ray on every configuration. Spark improved with distribution; Ray got worse with distribution.* 196s vs 259s single-node (Spark 24% faster), 117s vs 439s distributed (Spark 73% faster). This is the opposite of conventional “Ray is faster because Python-native” wisdom.




Figure 5: Total time across all four configurations. Spark goes 196s→117s with distribution (-40%); Ray goes 259s→439s with distribution (+70%).




Figure 6: Per-stage 2-node Spark vs Ray. Spark wins every stage; the gap is largest at ingestion and cleansing where Ray pays heavy network and serialization overhead.




Figure 7: Single-node vs distributed, framework-by-framework. Spark’s export drops 3.4× via parallel parquet writes. Ray’s every stage gets slower – LAN orchestration overhead exceeds parallelism gains at this scale.

5.2 Resource Utilization

Resource	Spark 2-node	Ray 2-node
Cores in use	8 / 8 (100%)	17 / 18 (94%)
M1 memory peak	2.80 GiB / 8 GiB (35%)	3.00 GiB / 8 GiB (37.5%)
M4 memory peak	4.00 GiB / 16 GiB (25%)	6.06 GiB / 16 GiB (37.9%)
M1 object store	–	919 MiB / 1.4 GiB (64%)
M4 object store	–	1.07 GiB / 1.4 GiB (76%)
Disk spill	negligible	4275 MiB

Table 4: Peak resource utilization (Spark UI / Ray Dashboard).

The 4275 MiB of disk spill on Ray is itself diagnostic: that volume of intermediate writes through pyarrow adds tens of seconds of latency that Spark, with its mature shuffle manager and on-heap Tungsten memory, simply does not pay.

6 UDF Deep-Dive

The Python UDF is the deciding architectural test. The premise: Spark’s JVM core has to ship every row across a JVM↔Python boundary, while Ray runs Python natively.

6.1 What the JVM Tax Looks Like

For a plain `@F.udf` in Spark, every row is: serialised Tungsten→Arrow, sent over a Unix socket to a forked Python worker, decoded into pandas, computed, re-encoded to Arrow, sent back, deserialised to Tungsten. Steps 1, 2, 5, 6, 7 are pure overhead. `@F.pandas_udf` (which we used) amortises this over ~10K rows per batch, dramatically reducing the cost. **Ray skips all seven steps – the function runs in the same Python process holding the pandas batch.**

6.2 Why the Theoretical Advantage Vanishes Here

In our distributed run, the Spark UDF stage took 34.5 s while Ray’s took 86.0 s. Three structural reasons:

- (1) **We chose `pandas_udf` for Spark.** It uses Arrow as the wire format – the same in-memory format Ray uses internally. We essentially gave Spark Ray’s wire format. Plain `@F.udf` would have been 5–10× slower.
- (2) **Catalyst optimizer fuses operators.** Spark fuses `filter → project → udf → filter` into a single pass. Ray Data treats each `map_batches` as an independent stage with materialised intermediates.
- (3) **The dataset is too small to amortise Ray’s overhead.** Ray shines at TB-scale workloads with deep Python dependencies (custom PyTorch inference per row, GPU-bound computation). 19.5M rows of arithmetic UDFs is Spark’s home turf.

6.3 Verdict on the UDF Hypothesis

The popular framing – “Ray wins because no JVM tax” – is true *in the limit*, but only when the Python computation cost dominates framework overhead. For typical ETL with light per-row Python, Spark + `pandas_udf` closes the gap entirely and then beats Ray on Catalyst, broadcast joins, the mature shuffle manager, and vectorised Tungsten readers. **Ray’s architectural advantage is real but situational.**

7 Pipeline Parity

Property	Spark	Ray
Initial rows	19,493,620	19,493,620
Final rows	18,271,535	18,272,001
Divergence	466 rows (0.0026%)	

Table 5: Pipeline parity check: 0.003% row divergence – attributable to non-deterministic ordering inside `dropDuplicates` when ties exist. **Parity passed.**

8 Performance Tuning Note

Spark optimisations:

- **`pandas_udf` instead of `@F.udf`.** Initial row-by-row UDF hung indefinitely; switching to `pandas_udf` (Arrow batches) reduced the UDF stage 10×.

- **F.broadcast() on the zones table.** Replicates 265 rows to every executor, eliminating the $O(N)$ shuffle cost on join.
- **Schema normalization at read.** Canonical schema with per-file casting handles NYC TLC type drift.
- **Reduced executor memory to 3 GiB.** An initial 6 GiB ask kept the app WAITING because no worker had that much RAM free.

Ray optimisations:

- **RAY_ENABLE_WINDOWS_OR_OSX_CLUSTER=1** – required for multi-node Ray on macOS.
- **-object-store-memory=1500000000** on each node – bounds the plasma store given M1 RAM constraints.
- **Driver-side write for export.** `transformed.write_parquet()` from a distributed dataset failed because workers tried to write to local relative paths. Writing from the driver after collection trades a small parallelism loss for filesystem simplicity.
- **Broadcast-style merge for zones.** Ray Data lacks a true distributed join; we materialise zones on the driver and `merge()` into each batch.

9 Challenges Encountered

NYC TLC schema drift. `VendorID`, `passenger_count`, `PULocationID` switch between INT32 and INT64; `airport_fee` sometimes capitalised. *Fix:* explicit canonical schema, per-file normalization, then union.

Username path mismatch across Macs. Spark records absolute paths in the execution plan; M4 received `/Users/zsres/...` which doesn't exist on his disk (his username is `nikeshkumarmandal`). *Fix:* symbolic link on M4 redirecting `/Users/zsres/Documents/spark-vs-ray` to the actual project location. In production, HDFS/S3/NFS removes this issue entirely.

Ray's macOS multi-node restriction. Ray refuses multi-node clusters on macOS by default. *Fix:* `RAY_ENABLE_WINDOWS_OR_OSX_CLUSTER=1`.

Setuptools 80+ removed pkg_resources. Ray 2.9.3's dashboard imports `pkg_resources._vendor.packaging` removed in `setuptools` ≥ 81 . *Fix:* pinned `setuptools==67.8.0`.

Disk pressure on M4. First Ray run aborted with `OutOfDiskError` during shuffle spill (M4 was at 99% disk usage). *Fix:* cleared ~ 40 GB on his Mac, capped object store at 1.5 GB.

Ray distributed write. `write_parquet()` from a distributed dataset failed because workers tried to write to relative paths on their local disks. *Fix:* write from driver after collection.

10 Winner and Recommendation

For an AI-First project: Ray. Ray Train, Ray Tune, and Ray Serve are first-class citizens of the Python ML ecosystem; data flows naturally from Ray Data $\rightarrow$ PyTorch $\rightarrow$ Ray Train. No JVM means heavyweight Python deps (HuggingFace, PyTorch, vector embeddings) compose effortlessly. For a project where data engineering is a thin prelude to a long ML training loop, the per-stage Spark advantage is dwarfed by integration cost.

For a BI-First project: Spark. For ETL, reporting, dashboards, traditional warehouse pipelines, Spark wins decisively – both in our benchmark and in industry adoption. Catalyst, mature shuffle manager, broad ecosystem (Hive, Delta, Iceberg, MLlib), SQL interface, operational maturity (the Spark UI is a generation ahead of Ray's dashboard).

Single-node on a laptop: *Spark wins*. 196s vs 259s. Catalyst’s whole-stage codegen and Tungsten memory beat Ray’s distributed object store at small scale, where Ray pays overhead with no parallelism dividend.

11 AI Disclosure

In accordance with the Code of Conduct for Fair Use of AI specified in the assignment, the following table documents all AI assistance used in this project.

Activity	AI Usage Status
Logic translation between PySpark and Ray Data syntax (e.g., translating <code>pandas_udf</code> patterns to <code>map_batches</code>)	Used Claude (Anthropic) per the assignment’s explicit “Logic Translation” rule. All translated code was reviewed, modified, and tested locally on the cluster. Allowed .
Debugging help (Spark schema-merge errors on NYC TLC parquet drift, <code>pkg_resources</code> setup-tools incompatibility, Ray’s macOS multi-node restriction, nginx-style path issues across Macs, <code>OutOfDiskError</code> on Ray spill phase)	Used AI to interpret error messages and evaluate possible workarounds. The actual fixes (e.g., explicit canonical schema, <code>setuptools==67.8.0</code> pin, <code>RAY_ENABLE_WINDOWS_OR_OSX_CLUSTER=1</code> , symlinks across user paths, capping object store memory) were chosen and applied manually. Allowed (debugging) .
LaTeX report drafting and structural proofreading	Used AI to structure this report and to match the formatting of my previous DA5402 submission. All technical content describes the actual implementation and the actual measurements. Allowed (proofreading) .
Cluster orchestration decisions (master/worker assignment, network topology, Spark executor sizing, Ray object-store sizing, <code>pandas_udf</code> vs <code>@F.udf</code> choice, broadcast-join strategy, schema-normalization design, drift-handling for parquet files)	Authored manually. These are the core engineering decisions the assignment is grading. Not AI .
Performance benchmark numbers (all timings in Section 5, all resource utilization figures, all parity statistics)	No AI involvement. Numbers are the raw output of the actual cluster runs, captured directly from the timings JSON files written by my pipeline scripts and verifiable against the Spark UI / Ray Dashboard screenshots in this report. The assignment’s “No Synthetic Benchmarks” rule was followed strictly. Not AI .

Table 6: AI Usage Disclosure – DA5402 Assignment 8

11.1 “Explaining the Why” Test

The following design choices can be defended without reference to AI:

- **Choosing M1 as master and M4 as worker** – Master coordination is lightweight; the heavier executor work belongs on the more powerful M4 (16 GB vs 8 GB).
- **`pandas_udf` over `@F.udf`** – The naive row-by-row UDF caused OOM and disk spills on the 8 GB

M1 in single-node testing. Vectorised batch transfer via Arrow is the standard Spark optimisation.

- **Broadcast join for the 265-row zones table** – A small dimension table replicated to every executor avoids the $O(N)$ shuffle cost of a hash join.
- **Explicit schema for parquet ingest** – Real-world data engineering hygiene; the NYC TLC files have type drift across months that Spark’s vectorised reader rejects strictly.
- **Driver-side write for Ray export** – Ray Data’s distributed `write_parquet` requires shared filesystem coordination not present in our 2-Mac LAN setup; collecting to driver before write is the correct trade-off.
- **Capping Ray’s object store at 1.5 GB** – Prevents the M1’s modest RAM from being saturated by the plasma store, at the cost of more frequent disk spills (which we measured at 4275 MiB).

12 Conclusion

This benchmark delivered an honest, reproducible comparison on real two-node hardware. The findings are at odds with conventional wisdom: **Spark beat Ray on every configuration, and Ray actively got worse when distributed across two Macs over Wi-Fi LAN.** Three reasons drove the result: (1) Spark’s `pandas_udf` captures most of Ray’s Python-native advantage by using Arrow for IPC; (2) Catalyst fuses operators that Ray Data treats as independent stages; (3) 19.5M rows is below the scale at which Ray’s actor model amortises its scheduling overhead.

The most valuable lesson is one benchmarks rarely surface: the popular “Ray for AI, Spark for BI” framing is correct, but for the wrong reason. Ray’s edge is not in raw processing speed – it is in *integration* with the Python ML stack. When you measure raw pipeline performance, Spark’s decade of maturity wins. When you measure end-to-end ML iteration speed, Ray’s lack of a JVM boundary wins. **For practitioners: if your workload is “data → data,” use Spark. If it’s “data → model → data,” use Ray. Don’t pick by hype; pick by the shape of the boundary you cross most often.**